

DevikaJainProject2

April 10, 2025

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.feature_selection import SelectKBest, f_classif

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.metrics import classification_report, accuracy_score
from imblearn.over_sampling import SMOTE
```

```
[2]: from google.colab import files
uploaded = files.upload()
```

<IPython.core.display.HTML object>

Saving HR-Employee-Attrition.csv to HR-Employee-Attrition.csv

```
[3]: df = pd.read_csv("HR-Employee-Attrition.csv")
df.head()
```

```
[3]:
```

	Age	Attrition	BusinessTravel	DailyRate	Department	\
0	41	Yes	Travel_Rarely	1102		Sales
1	49	No	Travel_Frequently	279	Research & Development	
2	37	Yes	Travel_Rarely	1373	Research & Development	
3	33	No	Travel_Frequently	1392	Research & Development	
4	27	No	Travel_Rarely	591	Research & Development	

	DistanceFromHome	Education	EducationField	EmployeeCount	EmployeeNumber	\
0		1	2 Life Sciences	1		1
1		8	1 Life Sciences	1		2
2		2	2 Other	1		4
3		3	4 Life Sciences	1		5

4	2	1	Medical	1	7
---	---	---	---------	---	---

	...	RelationshipSatisfaction	StandardHours	StockOptionLevel	\
0	...	1	80	0	
1	...	4	80	1	
2	...	2	80	0	
3	...	3	80	0	
4	...	4	80	1	

	TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance	YearsAtCompany	\
0	8	0	1	6	
1	10	3	3	10	
2	7	3	3	0	
3	8	3	3	8	
4	6	3	3	2	

	YearsInCurrentRole	YearsSinceLastPromotion	YearsWithCurrManager
0	4	0	5
1	7	1	7
2	0	0	0
3	7	3	0
4	2	2	2

[5 rows x 35 columns]

```
[4]: df.shape
      df.info()
      df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   1470 non-null   int64
1   Attrition                           1470 non-null   object
2   BusinessTravel                       1470 non-null   object
3   DailyRate                           1470 non-null   int64
4   Department                           1470 non-null   object
5   DistanceFromHome                     1470 non-null   int64
6   Education                             1470 non-null   int64
7   EducationField                       1470 non-null   object
8   EmployeeCount                        1470 non-null   int64
9   EmployeeNumber                       1470 non-null   int64
10  EnvironmentSatisfaction               1470 non-null   int64
11  Gender                               1470 non-null   object
12  HourlyRate                           1470 non-null   int64
```

13	JobInvolvement	1470 non-null	int64
14	JobLevel	1470 non-null	int64
15	JobRole	1470 non-null	object
16	JobSatisfaction	1470 non-null	int64
17	MaritalStatus	1470 non-null	object
18	MonthlyIncome	1470 non-null	int64
19	MonthlyRate	1470 non-null	int64
20	NumCompaniesWorked	1470 non-null	int64
21	Over18	1470 non-null	object
22	OverTime	1470 non-null	object
23	PercentSalaryHike	1470 non-null	int64
24	PerformanceRating	1470 non-null	int64
25	RelationshipSatisfaction	1470 non-null	int64
26	StandardHours	1470 non-null	int64
27	StockOptionLevel	1470 non-null	int64
28	TotalWorkingYears	1470 non-null	int64
29	TrainingTimesLastYear	1470 non-null	int64
30	WorkLifeBalance	1470 non-null	int64
31	YearsAtCompany	1470 non-null	int64
32	YearsInCurrentRole	1470 non-null	int64
33	YearsSinceLastPromotion	1470 non-null	int64
34	YearsWithCurrManager	1470 non-null	int64

dtypes: int64(26), object(9)

memory usage: 402.1+ KB

```
[4]:
```

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount \
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0
mean	36.923810	802.485714	9.192517	2.912925	1.0
std	9.135373	403.509100	8.106864	1.024165	0.0
min	18.000000	102.000000	1.000000	1.000000	1.0
25%	30.000000	465.000000	2.000000	2.000000	1.0
50%	36.000000	802.000000	7.000000	3.000000	1.0
75%	43.000000	1157.000000	14.000000	4.000000	1.0
max	60.000000	1499.000000	29.000000	5.000000	1.0

	EmployeeNumber	EnvironmentSatisfaction	HourlyRate	JobInvolvement \
count	1470.000000	1470.000000	1470.000000	1470.000000
mean	1024.865306	2.721769	65.891156	2.729932
std	602.024335	1.093082	20.329428	0.711561
min	1.000000	1.000000	30.000000	1.000000
25%	491.250000	2.000000	48.000000	2.000000
50%	1020.500000	3.000000	66.000000	3.000000
75%	1555.750000	4.000000	83.750000	3.000000
max	2068.000000	4.000000	100.000000	4.000000

	JobLevel ...	RelationshipSatisfaction	StandardHours \
count	1470.000000 ...	1470.000000	1470.0

mean	2.063946	...	2.712245	80.0
std	1.106940	...	1.081209	0.0
min	1.000000	...	1.000000	80.0
25%	1.000000	...	2.000000	80.0
50%	2.000000	...	3.000000	80.0
75%	3.000000	...	4.000000	80.0
max	5.000000	...	4.000000	80.0

	StockOptionLevel	TotalWorkingYears	TrainingTimesLastYear	\
count	1470.000000	1470.000000	1470.000000	
mean	0.793878	11.279592	2.799320	
std	0.852077	7.780782	1.289271	
min	0.000000	0.000000	0.000000	
25%	0.000000	6.000000	2.000000	
50%	1.000000	10.000000	3.000000	
75%	1.000000	15.000000	3.000000	
max	3.000000	40.000000	6.000000	

	WorkLifeBalance	YearsAtCompany	YearsInCurrentRole	\
count	1470.000000	1470.000000	1470.000000	
mean	2.761224	7.008163	4.229252	
std	0.706476	6.126525	3.623137	
min	1.000000	0.000000	0.000000	
25%	2.000000	3.000000	2.000000	
50%	3.000000	5.000000	3.000000	
75%	3.000000	9.000000	7.000000	
max	4.000000	40.000000	18.000000	

	YearsSinceLastPromotion	YearsWithCurrManager
count	1470.000000	1470.000000
mean	2.187755	4.123129
std	3.222430	3.568136
min	0.000000	0.000000
25%	0.000000	2.000000
50%	1.000000	3.000000
75%	3.000000	7.000000
max	15.000000	17.000000

[8 rows x 26 columns]

```
[5]: # Check for nulls and basic info
print(df.info())
print(df['Attrition'].value_counts())

# Visualize class imbalance
sns.countplot(data=df, x='Attrition')
plt.title("Attrition Distribution")
```

```
plt.show()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1470 entries, 0 to 1469
```

```
Data columns (total 35 columns):
```

#	Column	Non-Null Count	Dtype
0	Age	1470 non-null	int64
1	Attrition	1470 non-null	object
2	BusinessTravel	1470 non-null	object
3	DailyRate	1470 non-null	int64
4	Department	1470 non-null	object
5	DistanceFromHome	1470 non-null	int64
6	Education	1470 non-null	int64
7	EducationField	1470 non-null	object
8	EmployeeCount	1470 non-null	int64
9	EmployeeNumber	1470 non-null	int64
10	EnvironmentSatisfaction	1470 non-null	int64
11	Gender	1470 non-null	object
12	HourlyRate	1470 non-null	int64
13	JobInvolvement	1470 non-null	int64
14	JobLevel	1470 non-null	int64
15	JobRole	1470 non-null	object
16	JobSatisfaction	1470 non-null	int64
17	MaritalStatus	1470 non-null	object
18	MonthlyIncome	1470 non-null	int64
19	MonthlyRate	1470 non-null	int64
20	NumCompaniesWorked	1470 non-null	int64
21	Over18	1470 non-null	object
22	OverTime	1470 non-null	object
23	PercentSalaryHike	1470 non-null	int64
24	PerformanceRating	1470 non-null	int64
25	RelationshipSatisfaction	1470 non-null	int64
26	StandardHours	1470 non-null	int64
27	StockOptionLevel	1470 non-null	int64
28	TotalWorkingYears	1470 non-null	int64
29	TrainingTimesLastYear	1470 non-null	int64
30	WorkLifeBalance	1470 non-null	int64
31	YearsAtCompany	1470 non-null	int64
32	YearsInCurrentRole	1470 non-null	int64
33	YearsSinceLastPromotion	1470 non-null	int64
34	YearsWithCurrManager	1470 non-null	int64

```
dtypes: int64(26), object(9)
```


```
memory usage: 402.1+ KB
```

```
None
```

```
Attrition
```

```
No      1233
```


Yes 237
Name: count, dtype: int64

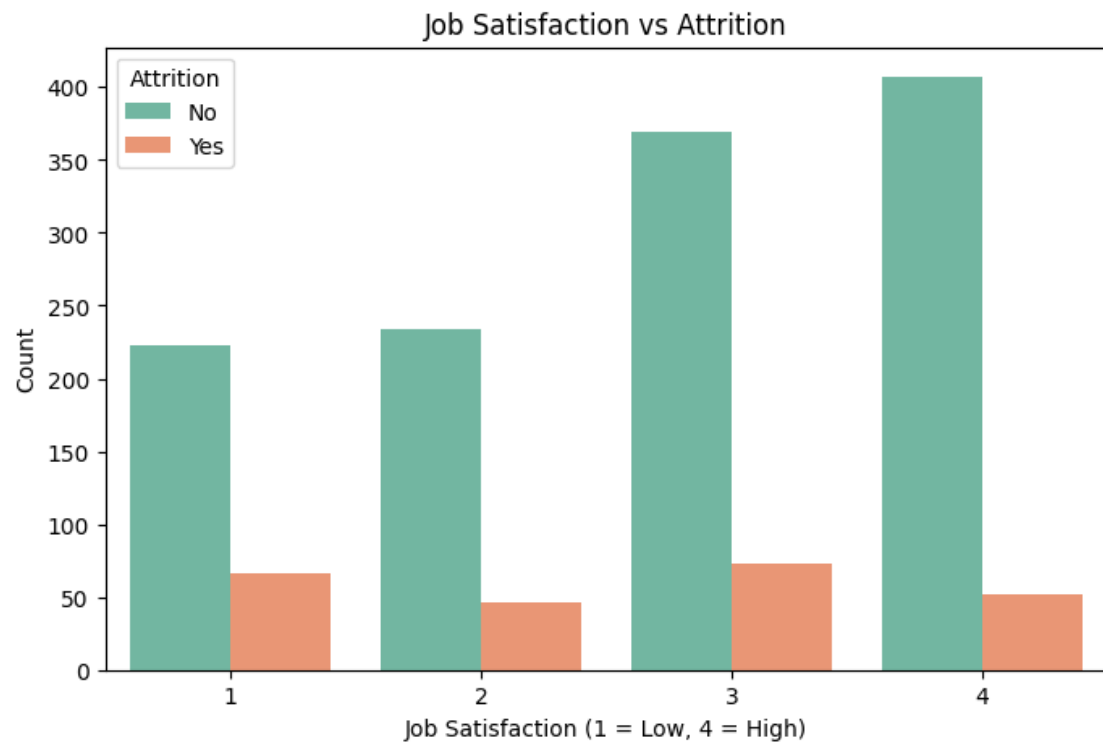
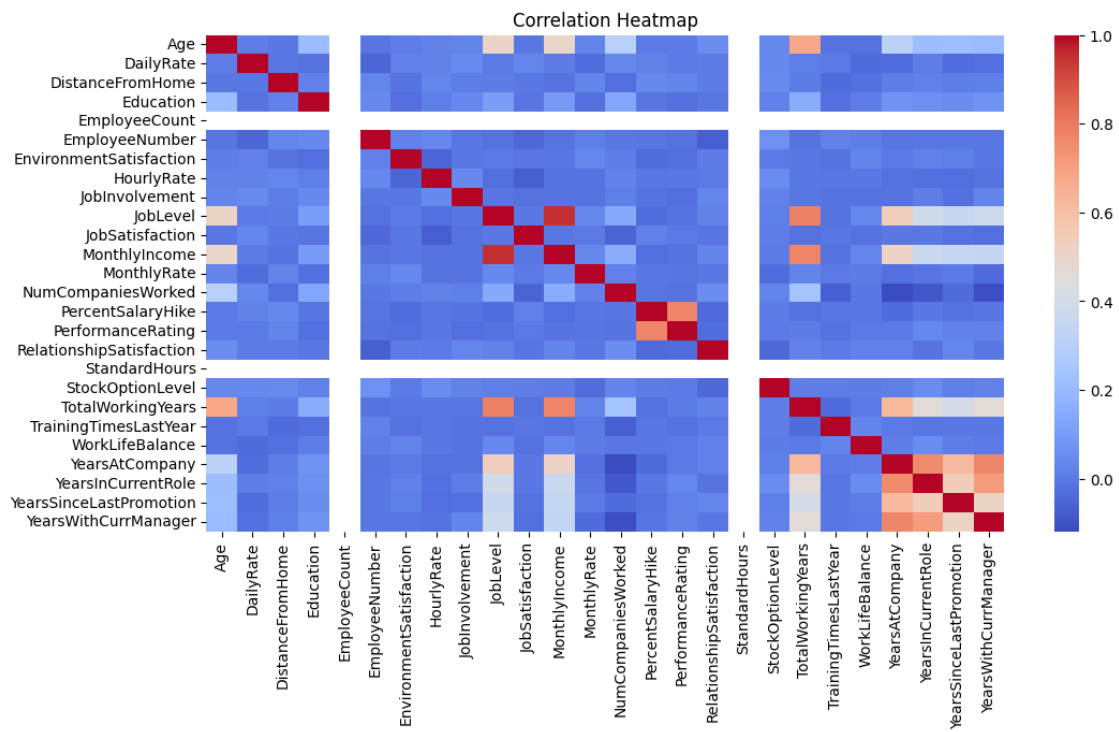


```
[6]: # Check target distribution
sns.countplot(data=df, x='Attrition')
plt.title("Attrition Count")
plt.show()
# This graph displays the distribution of the target variable - whether an
# employee has left the company (Attrition = Yes)
# or not (Attrition = No).
# It helps us understand the class imbalance in the dataset.
# If one class (e.g., "No") is significantly higher than the other, it
# indicates the need for balancing techniques like SMOTE.

# Correlation Heatmap
plt.figure(figsize=(12, 6))
sns.heatmap(df.corr(numeric_only=True), cmap='coolwarm', annot=False)
plt.title("Correlation Heatmap")
plt.show()
```

```
# This heatmap displays the correlation coefficients between all numerical
  ↳ features in the dataset.
# Values range from -1 to +1.
#   +1 indicates a strong positive correlation.
#   -1 indicates a strong negative correlation.
#   0 means no linear correlation.
# This helps us identify features that might be strongly related, which can
  ↳ guide feature selection or removal to avoid multicollinearity.
# Job Satisfaction Countplot split by Attrition
plt.figure(figsize=(8, 5))
sns.countplot(data=df, x='JobSatisfaction', hue='Attrition', palette='Set2')
plt.title("Job Satisfaction vs Attrition")
plt.xlabel("Job Satisfaction (1 = Low, 4 = High)")
plt.ylabel("Count")
plt.legend(title="Attrition")
plt.show()
# If attrition is higher at satisfaction level 1, it suggests low satisfaction
  ↳ is linked to higher attrition.
# If level 4 shows mostly "No", it implies higher satisfaction may retain
  ↳ employees better.
```






```
[7]: # Drop unnecessary columns
df = df.drop(['EmployeeCount', 'EmployeeNumber', 'Over18', 'StandardHours'],
             ↪axis=1)
# we dropped :
# EmployeeCount: This column has the same value for all employees(usually 1),so
             ↪it doesn't provide any useful information for the model.
# EmployeeNumber: It's a unique identifier for each employee, like an ID. It
             ↪creates noise.
# Over18: Like EmployeeCount, it's constant across the dataset, so it adds no
             ↪value
# StandardHours: All employees have the same standard hours (usually 40).No
             ↪variation = no influence on model prediction.
```

```
[8]: # Encode categorical variables
le = LabelEncoder()
for col in df.select_dtypes(include=['object']).columns:
    df[col] = le.fit_transform(df[col])
```

```
[9]: # Impute missing values (if any)
imputer = SimpleImputer(strategy='mean')
df = pd.DataFrame(imputer.fit_transform(df), columns=df.columns)
```

```
[10]: # Feature Scaling
scaler = StandardScaler()
# StandardScaler normalizes the features so they all have:
# Mean = 0
# Standard Deviation = 1
X_scaled = pd.DataFrame(scaler.fit_transform(df.drop('Attrition', axis=1)),
             ↪columns=df.drop('Attrition', axis=1).columns)
y = df['Attrition']
# since in ml:
# X = Features (input variables we use to make predictions)
# y = Target (what we want to predict)
# X_scaled = All employee-related attributes (like Age, JobSatisfaction,
             ↪Department, etc.)
# y = Whether the employee left the company or not → the Attrition column
             ↪(usually "Yes"/"No", or 0/1 after encoding)
```

```
[11]: # Show class distribution before SMOTE
print("Before SMOTE - Class distribution:")
print(y.value_counts())

# Handle Class Imbalance with SMOTE
sm = SMOTE(random_state=42)
X_resampled, y_resampled = sm.fit_resample(X_scaled, y)
```

```
# Show class distribution after SMOTE
print("\nAfter SMOTE - Class distribution:")
print(pd.Series(y_resampled).value_counts())
```

Before SMOTE - Class distribution:

Attrition

0.0 1233

1.0 237

Name: count, dtype: int64

After SMOTE - Class distribution:

Attrition

1.0 1233


0.0 1233

Name: count, dtype: int64

```
[12]: plt.figure(figsize=(10, 4))
```

```
# Before SMOTE
plt.subplot(1, 2, 1)
sns.countplot(x=y)
plt.title("Before SMOTE")
plt.xlabel("Attrition")

# SMOTE (Synthetic Minority Over-sampling Technique) creates synthetic samples
↳ of
# the minority class (e.g., employees who left) so that the dataset becomes
↳ balanced.
# Purpose: Fix imbalance so the model won't be biased toward the majority class.
# After SMOTE
plt.subplot(1, 2, 2)
sns.countplot(x=y_resampled)
plt.title("After SMOTE")
plt.xlabel("Attrition")
plt.tight_layout()
plt.show()
```



```
[13]: # Feature Selection
selector = SelectKBest(score_func=f_classif, k=10)
# SelectKBest is a feature selection technique.
# It selects the top k features that have the strongest relationship with the
  ↳ target variable (y_resampled, which is Attrition here).
# You set k=10, so it will keep only the 10 most important features.
X_selected = selector.fit_transform(X_resampled, y_resampled)
selected_features = X_scaled.columns[selector.get_support()]
print("Top features:", selected_features)
```

```
Top features: Index(['Age', 'JobLevel', 'MaritalStatus', 'MonthlyIncome',
'OverTime',
'StockOptionLevel', 'TotalWorkingYears', 'YearsAtCompany',
'YearsInCurrentRole', 'YearsWithCurrManager'],
dtype='object')
```

```
[14]: X_train, X_test, y_train, y_test = train_test_split(X_selected, y_resampled,
  ↳ test_size=0.2, random_state=42)

models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Random Forest": RandomForestClassifier(),
    "Support Vector Machine": SVC(probability = True)
}
model_accuracies = {}
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    model_accuracies[name] = acc
    print(f"--- {name} ---")
    print(f"{name} Accuracy: {acc:.4f}")
```

```
print("Classification Report:\n", classification_report(y_test, y_pred))
```

--- Logistic Regression ---

Logistic Regression Accuracy: 0.6943

Classification Report:

	precision	recall	f1-score	support
0.0	0.71	0.68	0.69	250
1.0	0.68	0.71	0.70	244
accuracy			0.69	494
macro avg	0.69	0.69	0.69	494
weighted avg	0.69	0.69	0.69	494

--- Random Forest ---

Random Forest Accuracy: 0.8664

Classification Report:

	precision	recall	f1-score	support
0.0	0.85	0.89	0.87	250
1.0	0.88	0.84	0.86	244
accuracy			0.87	494
macro avg	0.87	0.87	0.87	494
weighted avg	0.87	0.87	0.87	494

--- Support Vector Machine ---

Support Vector Machine Accuracy: 0.7429

Classification Report:

	precision	recall	f1-score	support
0.0	0.75	0.75	0.75	250
1.0	0.74	0.74	0.74	244
accuracy			0.74	494
macro avg	0.74	0.74	0.74	494
weighted avg	0.74	0.74	0.74	494

```
[15]: for name, model in models.items():
      scores = cross_val_score(model, X_selected, y_resampled, cv=5)
      print(f"{name} Average CV Score: {np.mean(scores):.4f}")
```

Logistic Regression Average CV Score: 0.7271

Random Forest Average CV Score: 0.8906

Support Vector Machine Average CV Score: 0.7851

```
[16]: import matplotlib.pyplot as plt
import seaborn as sns

# Plot accuracies
plt.figure(figsize=(8, 5))
bars = sns.barplot(
    x=list(model_accuracies.keys()),
    y=list(model_accuracies.values()),
    palette='viridis',
    legend=False
)

plt.ylabel("Accuracy")
plt.title("Model Accuracies")
plt.ylim(0, 1)
plt.xticks(rotation=15)

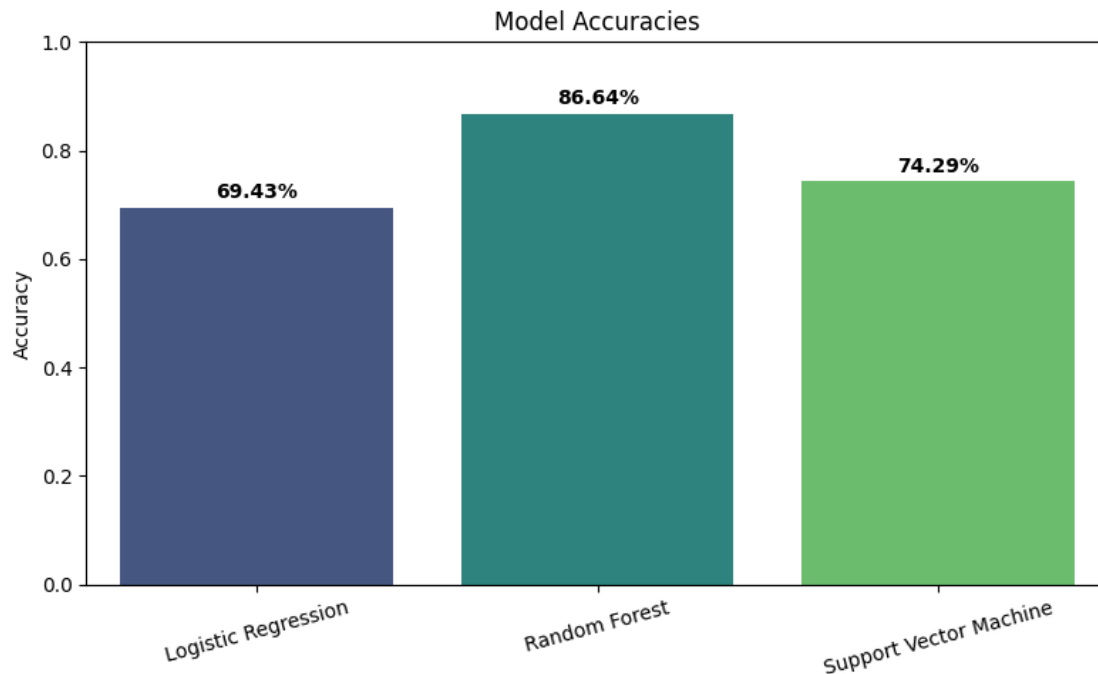
for bar in bars.patches:
    height = bar.get_height()
    plt.text(
        bar.get_x() + bar.get_width() / 2,
        height + 0.01,
        f"{height * 100:.2f}%",
        ha='center',
        va='bottom',
        fontsize=10,
        fontweight='bold'
    )

plt.tight_layout()
plt.show()
```

<ipython-input-16-3143671cc981>:6: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
bars = sns.barplot(
```



```
[17]: from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.feature_selection import SelectKBest, f_classif

# --- Feature Selection (already done) ---
selector = SelectKBest(score_func=f_classif, k=10)
X_selected = selector.fit_transform(X_resampled, y_resampled)
selected_features = X_scaled.columns[selector.get_support()]
print("Selected Top Features:", selected_features)

# --- Train-Test Split ---
X_train, X_test, y_train, y_test = train_test_split(X_selected, y_resampled,
    ↪ test_size=0.2, random_state=42)

# --- Hyperparameter Tuning for Random Forest ---
param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2],
    'bootstrap': [True, False]
}

rf = RandomForestClassifier(random_state=42)
```

```

grid_search = GridSearchCV(estimator=rf, param_grid=param_grid,
                           cv=5, n_jobs=-1, verbose=1, scoring='accuracy')

grid_search.fit(X_train, y_train)

print("Best Parameters Found:", grid_search.best_params_)
print("Best Accuracy Score:", grid_search.best_score_)

# --- Evaluate on Test Data ---
best_rf = grid_search.best_estimator_
y_pred = best_rf.predict(X_test)
print("Test Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))

```

Selected Top Features: Index(['Age', 'JobLevel', 'MaritalStatus',
'MonthlyIncome', 'OverTime',
'StockOptionLevel', 'TotalWorkingYears', 'YearsAtCompany',
'YearsInCurrentRole', 'YearsWithCurrManager'],
dtype='object')

Fitting 5 folds for each of 48 candidates, totalling 240 fits

Best Parameters Found: {'bootstrap': False, 'max_depth': None,
'min_samples_leaf': 2, 'min_samples_split': 2, 'n_estimators': 200}

Best Accuracy Score: 0.8965507935488016

Test Accuracy: 0.8805668016194332

Classification Report:

	precision	recall	f1-score	support
0.0	0.86	0.92	0.89	250
1.0	0.91	0.84	0.87	244
accuracy			0.88	494
macro avg	0.88	0.88	0.88	494
weighted avg	0.88	0.88	0.88	494

```

[18]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV, train_test_split
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt

# Split the data (assuming X_selected and y_resampled are already defined)
X_train, X_test, y_train, y_test = train_test_split(X_selected, y_resampled,
                                                    test_size=0.2, random_state=42)

# Untuned Random Forest
rf_default = RandomForestClassifier(random_state=42)
rf_default.fit(X_train, y_train)

```

```

y_pred_default = rf_default.predict(X_test)
acc_default = accuracy_score(y_test, y_pred_default)

# GridSearchCV for Tuned Random Forest
param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2],
    'bootstrap': [True, False]
}
grid_search = GridSearchCV(RandomForestClassifier(random_state=42), param_grid,
    cv=5, n_jobs=-1)
grid_search.fit(X_train, y_train)
best_rf = grid_search.best_estimator_
y_pred_tuned = best_rf.predict(X_test)
acc_tuned = accuracy_score(y_test, y_pred_tuned)

# Plotting comparison
plt.bar(['Untuned RF', 'Tuned RF'], [acc_default, acc_tuned], color=['skyblue',
    'orange'])
plt.title('Random Forest Accuracy Comparison')
plt.ylabel('Accuracy')
plt.ylim(0, 1)
plt.text(0, acc_default + 0.01, f'{acc_default:.4f}', ha='center')
plt.text(1, acc_tuned + 0.01, f'{acc_tuned:.4f}', ha='center')
plt.show()

```